

Alexander Taylor

A Purely Functional Approach to Graph Queries on a Database

Computer Science Tripos – Part II

St John's College

March 16, 2018

Proforma

Name: **Alexander Taylor**
College: **St John's College**
Project Title: **A Purely Functional Approach to Graph Queries on a**
Examination: **Computer Science Tripos – Part II, June 2018**
Word Count:
Project Originator: Adapted from a project proposal by Dr Eiko Yoneki
Supervisor: Dr Timothy Jones

Original Aims of the Project

The initial aim of the project was to produce a simple graph database system which wrapped over an SQL database, to expose an API treating the database as a purely functional datastructure. Operations were to be achieved in a monadic fashion, according to well defined DSL semantics. The success criteria were that a user could write a composable query with the DSL and receive correct results with appropriate error checking.

Work Completed

All success criteria were met and exceeded. I also implemented several extensions such as adding write functionality and a collection of bespoke backends using the LMDB key-value store. These new backends were evaluated by comparing their against the compiled SQL queries generated by the SQL backend.

Special Difficulties

None.

Declaration

I, Alexander Taylor of St John's College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed [signature]

Date [date]

Contents

1	Introduction	11
1.1	Progress	11
1.2	Motivations	11
1.3	Type safety	11
2	Preparation	13
2.1	Existing Graph Databases	13
2.1.1	Classes of Database	13
2.1.2	Schema	13
2.1.3	Mutability	13
2.1.4	Query Languages	14
2.2	Mutability	14
2.3	Query Language	14
2.3.1	DSL Syntax	14
2.3.2	Algebraic Datatypes	15
2.3.3	Typing	15
2.3.4	Semantics	15
2.3.5	Commands	15
2.4	Languages and Libraries Used	15
2.5	Software Engineering	16
2.5.1	Front End	16
2.5.2	Middle End	16
2.5.3	Back End	16
2.6	Scala Techniques	16
2.6.1	Type Enrichment	17
2.6.2	Implicit Parameters	17
2.6.3	Typeclass Pattern	17
3	Implementation	19
3.1	Note on Purity and Concurrency	19
3.2	Functional Programming Techniques	19
3.2.1	Monadic Compilation	19
3.2.2	Constrained Future Monad	19
3.2.3	Operation Monad	20
3.2.4	Local and Global State	20

3.3	Schema Implementation	21
3.3.1	Schema Hierarchy	21
3.3.2	DBObjects	21
3.3.3	Unerasure	21
3.3.4	Relations	21
3.3.5	SchemaDescription	22
3.3.6	Findables	22
3.4	Query ADT	22
3.5	Commands	22
3.6	DSL	23
3.7	Common Generic Algorithms	23
3.7.1	Simple Traversal	23
3.7.2	Full traversal	24
3.7.3	Pathfinding	24
3.7.4	Joins	24
3.8	Views and Commits	25
3.9	Memory Backend	25
3.9.1	Table Structure	25
3.9.2	Reads	26
3.9.3	Left Optimisation	26
3.9.4	Writes	26
3.9.5	Storage	26
3.9.6	Mutability	26
3.9.7	Pathfinding and fixed point traversal	26
3.10	PostgreSQL backend	27
3.10.1	Table Structure	27
3.10.2	Query Structure	27
3.10.3	Monadic Compilation	27
3.10.4	Writes	27
3.10.5	Mutability	28
3.10.6	Pathfinding and Fixed Point Traversal	28
3.10.7	Object Storage	28
3.11	LMDB Backends	28
3.11.1	Common	28
3.11.2	Original LMDB Implementation	30
3.11.3	Batched	30
3.11.4	Common Subexpression Elimination	30
3.11.5	Complex Common Subexpression Elimination	32
4	Evaluation	33
4.1	Unit Tests	33
4.2	Performance Tests	33
4.2.1	Introduction	33
4.2.2	Hardware	33

4.2.3	Datasets	33
4.2.4	Test Harness	34
4.2.5	Results	34
4.3	Semantics Proofs	34
4.3.1	Denotational == Operational	34
4.3.2	Typesafety	34
4.3.3	Join as a Monoid	34
5	Conclusion	35
	Bibliography	35

List of Figures

Acknowledgements

Acknowledgements

Chapter 1

Introduction

This project has been to build a new graph database management system. written a purely functional style in the language Scala. I have successfully built a modular, composable read oriented DBMS which has a range of commands and a domain specific read query language.

1.1 Progress

I have achieved all of my initial success criteria, written several extensions (Writes, an extra backend, and optimisations to the backend), and achieved performance comparable to and sometimes exceeding that of postgresQL across a range of benchmarks.

1.2 Motivations

Graph databases becoming of more interest recently, as the relational model of data is found not to match certain datasets of interest, such as social networks.

Similarly, purely functional, strongly typed programming is also progressing from a research interest to a serious paradigm used in industry. This is especially the case in fields where the ability to rapidly produce high quality code which has a high probability of correctness is vital, such as in automated trading. In such cases, an advanced type system can be used to make guarantees about the behaviour of the program at runtime and enhances the ability to find bugs at compile time rather than in the field, while the purely functional paradigm allows code to be written far more expressively and concisely than in lower level languages.

There exist several attempts to make access to relational databases purely functional or monadic in nature, but there is very little application to graph databases.

<https://hackage.haskell.org/package/haskelldb> <http://www.bcs.org/upload/pdf/ewic'dp95'paper14>.

1.3 Type safety

When we talk about type safety, we mean to say that a correctly typed program will not produce run-time errors of a given class. Sources of a lack of type safety in typical

programs include exceptions (These are inherently not type safe in scala, since exceptions are unchecked), the usage of primitive data types where specific values are expected, such as passing strings to function without wrapping or tagging them in more specific types, and unmarshalling data from a type-erased source. If we dont know whether were expecting a value to be an integer or a string and we try to read it as an integer, there is a chance to throw an exception. By this definition, accessing an external, imperative, database from scala is not type safe. The database will have some collection of error modes that often manifest as exceptions, and the database system typically stores data in a type erased form that needs to be unmarshalled. Finally, queries to external database systems are typically constructed as a primitive string that cannot be checked automatically at compile time. This project addresses these issues by using carefully constructed result container types to correctly handle error cases, including arbitrary exceptions using a pattern match-able type hierarchy, a DSL that is checked at compile time by the scala type checker, and typeclass based marshalling.

Chapter 2

Preparation

2.1 Existing Graph Databases

2.1.1 Classes of Database

Existing graph databases typically fall into two classes: property graphs and edge labelled graphs. <https://arxiv.org/pdf/1610.06264.pdf> . Edge labelled graphs typically store only a label on any given edge, whereas property graphs can store more attributes on edges, hence being closer to close to relational databases. Edge labelled graphs can be more succinctly modelled mathematically, since they can be represented purely using mathematical relations. Edge labelled graphs also lead to a cleaner syntax (we shall cover this later). Finally, any data expressible using a property graph can be represented using an edge labelled graph by introducing additional nodes to hold attributes.

Insert and acknowledge edge-labelled vs property graph examples in GDrive

2.1.2 Schema

Current graph database systems often do not have rigid schema <https://arxiv.org/pdf/1602.00503.pdf>, introduction. Instead they use dynamic schema. This requires additional typing checks at runtime. This also fails to fit into our definition of type safety, since we have no guarantees about any objects extracted from the database. Instead, Ive opted for a rigid schema which is built using scala objects and hence is checked by the scala compiler at compile time.

2.1.3 Mutability

Current graph databases tend to be mutable. In order to express concepts that require immutability, such as time-dependent data, the user has impose extra constraints, such as time stamping nodes and relations. A lack of immutability also complicates concurrent semantics and implementation of ACID transactions. The introduction of immutable data structures gives us these effectively for free.

2.1.4 Query Languages

Typically, database systems have their own query language, such as cypher for neo4j and gremlin. These are typically used by generating a query string and submitting it to the database engine. This is inherently un-typesafe. The job of ensuring that absolutely every query that a given program could generate is valid is intractable to undecidable for non trivial programs. By encoding our query language in the host language and aligning the type system to be checked by the host compiler, we can get much stronger guarantees about program correctness. [links to cypher and gremlin](#)

2.2 Mutability

In order to implement immutability, we need a way of creating an immutable snapshot of the database at a given point in time. This is done using a system of views. A view is an immutable view of the database. When we read from the database, we read from a particular given view. When we write to a particular view of the database, we copy the view, update it, and return the new view. Reads and writes now never interfere with each other. [Read/Write diagram](#)

2.3 Query Language

2.3.1 DSL Syntax

I spent some time iterating over a DSL syntax and how I wanted queries to look and be structured. My goals were to have a highly composable and expressive, yet small language. At first I experimented with a neo4j style language, with pattern matching syntax, an example of which is below. The query takes a number of fields to fill in, in this case a pair of actors, and finds all valid ways to fill them in subject to the graph constraints. [Ugly syntax figure](#) However, the scala type inference is fairly localised so would need some type parameterisations filled out, yielding a more accurate syntax below. [Even uglier syntax](#) This is fairly cluttered syntax and difficult to read. It would also have been complicated by a desire to implement property graph features. Furthermore, we would not be able to easily and safely replicate Neo4js ability to extract an arbitrary number of objects of different types from a path. Doing so would require use of structures such as heterogeneously typed lists.

As a result, I settled upon an edge-labelled-relation oriented approach which neatens the above query significantly. [reword above; neat syntax](#)

This is much more readable and concise. With this syntax, and the assumption that [ActsIn](#) relates [Actors](#) to [Movies](#), Scalas typesystem can infer that [coactor](#) relates [Actors](#) to other [Actors](#). There is also syntax for repetitions, intersections and unions. [Examples in the appendix](#)

2.3.2 Algebraic Datatypes

These DSL queries should correspond to an intermediate representation. I have picked a set of expression constructors. These fall into two kinds: `FindPair` and `FindSingle`, indicating whether they return pairs or individual values. *Insert latex of the ADTs here*

We also define a view of the database and schema, as well as findables

Import view + schema semantic definitions here; reword if necessary

2.3.3 Typing

In order to ensure the correctness of a given query, there is a type system which checks the correctness of queries. *Latex for typing; Where do I mention that this is encoded into scala?*

2.3.4 Semantics

To fully define the query language, we need to define the semantics. I have defined two collections of semantics: operational style and denotational style. *import the latex for these; Proof?* I have also proved in the appendix the correspondence of these two sets of semantics.

2.3.5 Commands

In addition to the query construction language, there are five commands which make use of the queries. *define and semantics*

2.4 Languages and Libraries Used

Scala An object-oriented, functional JVM languages interoperable with Java. I chose Scala for its familiarity, advanced type system, ability to use Java libraries, and ease of deployment across many systems. <https://github.com/tperrigo/scala-type-system/wiki/Scala-Type-System-Overview>

Scala Build Tool A package manager and build tool for Scala. <https://www.scala-sbt.org/>

PostgreSQL An open source, multi platform, performant SQL implementation. <https://www.postgresql.org/>

LMDB An open source, highly optimised key-value datastore. <https://symas.com/lmdb/>

Scalaz A library for scala providing typeclasses and syntax to aid advanced functional programming. <https://github.com/scalaz/scalaz>

JDBC Javas standard library support for SQL connections. Used to interact with a postgres database. <http://www.oracle.com/technetwork/java/javase/jdbc/index.html>

LMDBJava A JNI library allowing access to the LMDB datastore. <https://github.com/lmdbjava/lmdbjava>

Junit A unit testing library for JVM languages. <https://junit.org/junit5/>

SLF4J A logging library for JVM languages <https://www.slf4j.org/>

Spray-JSON A JSON library for Scala, used to generate benchmarking datasets. <https://github.com/spray/spray-json>

Python Used to generate test datasets. <https://www.python.org/>

2.5 Software Engineering

The system was designed to be as modular and re-implementable as possible.

[Figure from Google drive: the architecture map](#)

2.5.1 Front End

The front end mostly consists of interfaces and syntax for building queries and the schema for the database.

2.5.2 Middle End

The front-ends DSL translates into the underlying typed ADT. This then has its compile time type information erased to become an un-typed query ADT, which is easier and cleaner to interpret.

2.5.3 Back End

Finally, there are three interfaces that a backend system must implement: DBBackend, DBInstance and DBExecutor. These specify a root backend object that opens an instance which represents a database connection. Each DBInstance has an executor that allows us to execute commands

2.6 Scala Techniques

I have made use of several advanced programming techniques which are specific to the Scala language.

2.6.1 Type Enrichment

The type enrichment feature of the scala language allows retroactive addition of methods to previously defined types. Type enrichment is performed by creating an implicit class, taking a single underlying value of some given type. The methods defined in the implicit class can now be called as if they were methods of the underlying class, provided the implicit class is in scope. **Type enrichment example from G Drive** This allows us to add syntax to types where a small number of core methods have been defined.

2.6.2 Implicit Parameters

Another advanced feature of scala that I have used is that of implicit parameters to functions and class constructors. Values such as vals, functions, and classes can be declared with a implicit tag. Functions and classes can declare additional parameters as implicit. When these functions are called, the implicit parameter can be omitted if and only if there is an unambiguous implicit value of the correct type in scope. **Implicit parameters example 1** This is typically used for purposes such as passing around values that are typically defined once and used many times in a program, such as an ExecutionContext, or a logging framework instance.

Functions can also be implicit and take implicit values **Implicit function example, caption: When an implicit value of type A is available, then an implicit value of type B is also available.**

This can be made more interesting when we include parameterised generic types. This allows us to get the compiler to do work in the manner of an automated theorem prover (by the curry-howard correspondence) at compile time.

Implicit parameters example 2

2.6.3 Typeclass Pattern

A further combination of these two patterns is the typeclass pattern. We define a typeclass for a type by defining as methods on a trait the operations we want on the type. We can then define implicit objects which work as type class instances for the types we want. We can also use implicit functions to generate typeclass instances in a manner similar to a proof tree. **example**

We can then use the type enrichment feature to add methods to values of a type that is a member of the typeclass.

Monoid typeclass example. Caption: An example of combining the three techniques above to define a Monoid Typeclass, ways to construct instances of the monoid typeclass and a syntax object which enriches types for which monoid properties can be proven.

It is clear that these patterns allow for very expressive structures and abstractions to be built in scala. I use these frequently within the project to neaten code and to achieve type safety.

Chapter 3

Implementation

3.1 Note on Purity and Concurrency

All of the backends aim to preserve the global immutability of the database. The immutable semantics of the database also mean that queries generally don't interfere with each other. This means that we can avoid keeping locks or creating large transactions. Hence, the backends don't require much work to maintain correct concurrency.

3.2 Functional Programming Techniques

3.2.1 Monadic Compilation

At several points in building a backend, it becomes necessary to transform one algebraic type to another. This is typically done by walking over a tree, whilst keeping some mutable state representing parts of the tree that are relevant. One example is converting an intermediate representation to SQL output code. Here, we may want to extract common subexpressions into a dictionary, or pick out all the database tables that need to be accessed by the query. This can be encoded by folding a State Monad instance over the tree.

The state monad is an abstraction over functions that chain an immutable state through successive computations.

[State monad definition or link to it](#)

To define a monadic compiler, we define a recursive function which chains state monad objects together. [State monad compiler example](#)

3.2.2 Constrained Future Monad

Part of the goal of type safety is the recovery of error cases. Typically, this would be done in a JVM program through the use of exceptions. However the presence of unchecked exceptions on the Java platform makes it difficult to ensure that all error cases are accounted for. A more functional approach is the use of the exception monad. <http://homepages.inf.ed.ac.uk/wadler/papers/marktoberdorf/baastad.pdf> In scala, this manifests as the built in Try monad and Scalaz's `\|` (`Either`) monad. `Try[A]` has two case

classes: `Failure(e: Throwable)` and `Success(a: A)`, while `E \/A` has the case classes `-/(e: E)` and `\/(a: A)` (`Left` and `Right`). Since `Try`'s failure case is the unsealed `Throwable` trait, we don't really have a way to ensure we have handled all error cases at compile time, whereas `Either` has a parametrised error type, which can be a sealed type hierarchy, which is then checked by the scala compiler at runtime. For example, consider the simple interpreter below. All error cases are proved to be handled by the typesystem. **Error interpreter example**

Typically, we're dealing with cases which might take a significant amount of time to return. So rather than using a simple `Either` monad, we lift it into an asynchrony monad. There are several options to choose from for an asynchrony monad. I chose the built in `Future` over more exotic alternatives, such as the scalaz `Task`, since it is relatively widely used, and I have some familiarity with it from past projects. `Futures` also capture thrown unchecked exceptions, which makes handling them a little easier. Hence we're interested in passing around `Future[E \/A]` around, for a sealed type hierarchy `E`. There is also the issue of the java libraries used (for SQL and LMDB access) throwing exceptions, and unexpected exceptions turning up in code. Fortunately, the `Future` container catches these, acting like an asynchronous `Try[E \/A]`. This causes issues as we don't have the sealed trait property of errors as we have above. To solve this, I introduced the `ConstrainedFuture[E, A]` monad, which has the requirement that the error case type parameter `E` implements the `HasRecovery` typeclass for converting any `Throwable` to an `E`.

Hashrecovery typeclass

The underlying future is kept private, and can only be accessed via the `run` method, which calls the `recover` method on any errors (tail recursively, so any exceptions thrown during execution are also handled). By this construction **appendix** we ensure all non-fatal error cases are contained in a type-safe way.

3.2.3 Operation Monad

As stated, **Can I link to it?** typical database operations take a view as a parameter, inspect the view, return a value and may also insert a new view. This requires interplay between the `ConstrainedFuture` (To handle failure and asynchrony) and `State` (to chain together updates to the view representing current state) monads. Hence we use the `Operation[E, A]` monad, which wraps a function `(ViewId => ConstrainedFuture[E, (ViewId, A)])` in a similar way to how `State` monad chains together functions `S => (S, A)`. Each of the commands on a database yields an operation.

3.2.4 Local and Global State

Although it would be preferable to only use purely functional folds, maps, and immutable data structures everywhere within the project, for certain, high frequency, performance critical tasks, using purely immutable structures slows us down. Recursive functions (though my functional style does not make heavy use of them anyway) tend to use more stack space than the JVM has available in many situations. Hence for tasks such as retrieving values for result sets, pathfinding across large relations, and building indices, I

have opted to make use of mutable data structures locally, using builders for collections such as sets. This leads to a dramatic increase in speed, especially for when large numbers of elements are added sequentially [Source?](#). In these cases, the mutable state never leaks out of the functions that make use of the mutability.

3.3 Schema Implementation

One of the goals of the project was to allow close to arbitrary user objects (assuming that they are finite) to be inserted and retrieved from the database. This was achieved using the typeclass pattern.

3.3.1 Schema Hierarchy

In order to work with the database, a class needs to have an instance of the [SchemaObject](#) typeclass. [definition of schema object](#)

[SchemaObject](#) is sealed, so can only be implemented by implementing one of its subclasses. Currently, for the sake of simplicity, there are five: [SchemaObject0](#) .. [SchemaObject4](#), with the number indicating the number of underlying database fields required. These are implemented by effectively defining marshalling functions from the type to and from tuples of [Storable](#) (read: primitive) types.

[To go in the appendix: The definition of a schema object subclass](#)

3.3.2 DBObjects

To store objects in the database under various backends, we need to have a type erased (at compile time, but not runtime) version of the objects. Once we have a primitive representation of an object from the [SchemaObject](#), we can fully unerase it by converting it to a [DBObject](#). A [DBObject](#) is simply a collection of type tagged fields ([DBCCell](#)). These can now be easily inserted or retrieved from a database. [DBObject and cell definitions](#)

3.3.3 Unerasure

In order to correctly retrieve values from the database, we need to be able to undo the erasure process. This can be done using the unmarshalling methods derived from the [SchemaObject](#) for an object type.

3.3.4 Relations

In order for type checking to work, relations need to have type parameters for the object types they link. Hence, to define relations in the schema, the user needs to define objects that extend the relation interface.

[Relation definition](#)

3.3.5 SchemaDescription

In order to build the database structures the backends need a definitive collection of schema to include. This is done using a [SchemaDescription](#) object, which simply holds a collection of [SchemaObjects](#) and [Relations](#) which need to be used by the database.

3.3.6 Findables

For the sake of simplicity, I have only implemented findables which test if fields of objects match particular values. This allows us to look for specific objects or match particular fields. This also makes indexing easier to implement.

3.4 Query ADT

This feels like it needs rewriting The ADTs described in the preparation section [link](#) are implemented in scala each by a pair of ADTs: a typed and type-erased equivalent for each. The typed ADTs are parameterised by the object types which they lookup. Due to this typing and scalas type inference, I have encoded the type of the typed ADTs such that the scala compiler checks the types and does inference for us, according to the type rules of the query language. **Appendix: Type definitions of ADT** The only rules that cannot be checked at compile time are whether the schema description contains the relation in instances of the ([Rel](#)) rule, the type [A](#) for the ([Id](#)) rule, and the ([Find](#)) rule, as we cannot predict the contents of the SchemaDescription at compile time without dependent types. These typed ADTs are erased, with respect to a schema, into their unsafe equivalents in order to be executed. If an AST node makes reference to a non-existent table or relation, a runtime error is created in the [Either](#) return type. The constructors of the type-erased ADT nodes are private to the enclosing package, meaning that they can only be created by erasing a typed ADT.

3.5 Commands

As specified in the previous chapter, each backend needs to implement five commands:

- [find\(S\)](#)
- [findPairs\(P\)](#)
- [ShortestPath\(start, end, P\)](#)
- [allShortestPaths\(start, P\)](#)
- [insert\(relations\).](#)

Each command should return an [Operation](#) of the correct type.

3.6 DSL

The DSL mostly consists of syntactic sugar to make queries easier to read. It is implemented using the type enrichment pattern. Both [Relation](#) and [FindPair](#) implement the trait (interface) [FindSingleAble](#), so we can use type enrichment to write new DSL operators. The main thing to note in the DSL is the use of arrows to chain relations together. (See [RelationSyntax.scala](#) in the appendix for examples of DSL). [Put in appendix](#)

3.7 Common Generic Algorithms

During construction of the database, several common patterns of problems emerged with slight differences. Hence, I have written relatively optimised generic versions of these algorithms such that different backends can make use of them regardless of the underlying types. These algorithms are found in [core.utils.algorithms](#)

3.7.1 Simple Traversal

The first set of generic algorithms to look at are the [SimpleFixedPointTraversal](#) algorithms. These compute the repetitions of a function for execution of the [FixedPoint](#), [Upto](#), and [Exactly](#) expressions of the ADT. They are labelled as Simple because they do the search from a single root. They carry out search mutably, and convert their output to an immutable set upon returning.

Exactly The simplest algorithm is for computing [Exactly](#) query nodes. Here, we simply expand a fringe set of values outwards, by applying the search step to every value in the fringe to get the new fringe. We also memoise the search step function in the case of repetitions. After the required number of repetitions, the remaining fringe is returned. [\(Diagram showing expanding fringe\)](#)

Upto The next algorithm is for computing [Upto](#). This is computed in a similar way by flat-mapping the functions over the fringe repeatedly to calculate an expanding fringe. The major differences here are that we keep an accumulator of all the found values. When a new fringe is calculated, we subtract the accumulator set from it to reduce the number of nodes that need to be searched to those that have not yet been searched, and then union the remaining fringe with the accumulator to get the new accumulator. After the required number of repetitions, the accumulator is returned. [diagram](#)

FixedPoint The final algorithm is to calculate [FixedPoint](#). This works slightly differently. As before, we keep an accumulator of reached nodes, but unlike before, the generation number of each node is not important, only that a node is reachable is important. Hence, the fringe is a queue rather than a set, and we iterate until the fringe is empty. In each iteration, we pop off the top value of the fringe queue, compute all

immediately reachable nodes. This reachable set is diffed (define diff) with the accumulator to find the newly reached nodes. These are now added to the fringe queue and the accumulator. When the fringe is empty, we return the accumulator.

3.7.2 Full traversal

The next set of algorithms build on the `SimpleFixedPoint` algorithms to return not just those nodes reachable from a single root, but the set of all reachable pairs with the left hand pair derived from a root set (using the left-hand optimisation). As such, the algorithms need to do some more work to reconstruct which nodes are reachable from each root, while still eliminating redundancies.

Exactly The first such algorithm is for computing `Exactly`. This is similar to the original version, except we now store a fringe for each root in a map of `Root => Set[Node]`. We also memoise the search function in a Map to avoid computing it redundantly. In each iteration, we simply expand the fringe for each root as before by mapping the fringe expansion loop body over the values of the fringe map. [Diagram](#)

Upto The next algorithm is to compute `Upto`. This is again done like before, but with a map of root to accumulator set as the accumulator. As with `Exactly`, the fringes of each root are expanded simultaneously, sharing redundant results via the memo, while the accumulators are unioned with the fringe of the appropriate root with each iteration. [Diagram](#)

FixedPoint Finally, the `FixedPoint` take a departure from the parallel implementations above. The reachable set of each node is calculated sequentially using a similar algorithm to above. However the memo now contains all nodes reachable from previously processed roots, allowing for fast convergence of dense graphs.

3.7.3 Pathfinding

The `ShortestPath` and `AllShortestPaths` commands require us to search a subgraph generated by a relation. Since all edges have unit weight, the pathfinding algorithm reduces to breadth-first-search, which I have implemented in an imperative format while wrapping up error cases in an `Either`. These functions take as a parameter the search step `A => E Set[A]` which represents the edges going out of a node. [Is this needed/clear?](#) There is room for optimisation here, since most backends focus on returning sets of pairs rather than just the right hand side.

3.7.4 Joins

The problem of joining two sets of pairs based on shared intermediate values is a requirement for any backend. [Latex definition of a join](#) I have implemented a simple hash join (<http://www.csd.uoc.gr/hy460/pdf/p63-mishra.pdf>). This operates by first assuming that the size of distinct leftmost elements in the right set is smaller (known to be a

subset of, due to the left-optimisation of backends) than the rightmost elements of the left set. We then build a mutable map to index leftmost values to rightmost values. We then traverse the left set and for each pair generate all new joined pairs, using a `explain that flatmap == bind?` flatmap to collect the values into one set. An issue with this approach is that we have to build the index upon each join call, an issue that is addressed later.

`Join algorithm code?`

3.8 Views and Commits

In the memory backend, as will be explained later, views are easy implement as a map of `ViewId` to `MemoryTree` and simply updating the `MemoryTree`, allowing Scalas immutable collections to handle sharing of data in an efficient way. In the other backends, backed by non-functional technologies, we need other methods of sharing and inserting to immutable structures. A first observation is that with the operation monad model, each view only has one direct predecessor, forming a tree where were only interested in the path back to the root from a given node. From this insight, we can think of each new view only needing to store a difference against its parent. Since the current design of the database only allows the addition of values and not deletion, this difference is only going to be positive. Hence, we can store all the added values between two views in a container called a `Commit`. As an optimisation, we can disregard the parent view, and simply store each view as a collection of its commits. This would also allow us to implement deletion if desired. This would be done by removing commits in a view that are to be modified, and replacing them with a new commit containing the contents of those commits, excluding the deleted values.

3.9 Memory Backend

The first backend that I have implemented is a simple, naive memory-based backend. This backend follows the denotational semantics, and makes very few attempts to improve performance. This backend serves as a test-bench backend, used to create unit tests to test other backends. It also allowed me to practice implementation of type erasure and unerasure according to the schema in a controllable environment (that is, without interference from other languages as libraries as in the SQL and LMDB backends.)

3.9.1 Table Structure

A memory instance stores a concurrent map of `ViewId` to `MemoryTree`, which itself is a map of `TableName` (derived from the `SchemaDescription`) to `MemoryTable`. There is a `MemoryTable` for each object class in the `SchemaDescription`. A `MemoryTable` provides lookups using maps for `DBObjects` and `Findables`, in the form of an index to set of objects for each column value. These lookups return objects containing a `DBObject` and the outgoing and incoming relations for the object, indexed by `RelationName`. `Diagram of View =i Tree(TableName =i Table(Index =i Object))`

3.9.2 Reads

Reads occur by simply walking over the ADTs recursively, following the denotational semantics [Code is small enough to place in an appendix](#). The only real departures from the denotational semantics are the Left-optimisation and use of the generic fixed point algorithms to compute [Exactly\(n, P\)](#), [Upto\(n, P\)](#), and [FixedPoint\(P\)](#). Results are also computed in the [Either](#) monad to allow for error checking (for cases such as missing tables).

3.9.3 Left Optimisation

One of the few optimisations here is the Left optimisation. When we compute the result set of a [FindPair](#), we pass in the subset of left hand side variables we want to compute from. This mostly has an effect when we compute joins ([Chains](#)). Consider joining a query of maybe a few dozen unique rightmost values to a large query with several million unique leftmost values. The pairs of the right relation only need to be joined if their leftmost value is in the set of rightmost values of the left relation. Hence it makes sense to pre-limit our search to pairs with leftmost values in the right most value. ([Diagram showing redundant calculations in the join](#)). This pattern also makes an appearance in the original LMDB implementation.

3.9.4 Writes

Thanks to Scalas immutable collections library, updating the database is fairly easy. When inserting a collection of relations, we simply add, relation by relation, each object in the relation, if it doesn't exist, and update the outgoing and incoming relation map of each object. This update creates new immutable object tables and a new memory tree. This is stored to the map of [ViewId](#) to [MemoryTree](#) as a new view.

3.9.5 Storage

Objects are stored as [DBObject](#)s in wrapper [MemoryObject](#)s. [MemoryObject](#)s are hashed and compared by their [DBObject](#)s. [Memory object diagram](#)

3.9.6 Mutability

The only mutability in the Memory implementation is for the views counter and the views lookup table. Access is kept transactional using a lock.

3.9.7 Pathfinding and fixed point traversal

The pathfinding and fixed point traversal methods are simply implemented by calls to the appropriate generic algorithms. One redundancy is that the search functions passed calculate all linked pairs, rather than optimising only to look at right hand sides of relations. ([more detail?](#))

3.10 PostgreSQL backend

Upon completion of the initial memory backend, I started work on a PostgreSQL based backend. This compiles the ADT intermediate representation into SQL queries that are then executed by a Postgres database.

3.10.1 Table Structure

The construction of the underlying database uses several SQL tables. These can be partitioned into a set of control tables that will be present in all database instances and a set of schema defined tables.

Control Tables [Insert the table from GDrive](#) Note: [ViewId](#) is a foreign key into the [views registry](#), [CommitId](#) is a foreign key into [Commits Registry](#). The [Dummy](#) column is needed to fix postgresql syntax when we try to get the next key for a table with one column.

Schema defined Tables [Insert the table from GDrive](#) Note: [Object-](#), [Left-](#), and [Right-](#)Ids are foreign keys to the [ObjectTables](#) [ObjectId](#) column.

3.10.2 Query Structure

3.10.3 Monadic Compilation

In order to generate these SQL queries, we need to convert the ADT query to a lower level intermediate representation [In appendix](#), which maps one-to-one to SQL. While doing this, we need to gather and rename all of the relation and auxiliary tables that we extract from so that we can form the CTE queries, we also need to find repeated queries to hoist out. In order to do this, we use the monadic compiler pattern described above. The compiler state is shown below [name the figure](#). When the compilation is done, depending on the context of the command, we append different extraction queries. For find pairs we need to extract the fields of both the left hand side and the right hand side objects, whereas for pathfinding operations, we only need to extract the [ObjectIds](#) along the path as opposed to whole objects. [Compilation context figure from GDrive](#)

3.10.4 Writes

There are several steps in the implementation of writes, though I have not expended a great deal of effort making them fast. A significant part of this is the insert or get SQL query, which looks up an object in the relevant table, returning its [ObjectId](#) if it exists and creating the object and returning the new ID if it does not.

[insertOrGet query from GDrive](#)

We create a new view and commit, then we run a memoised [InsertOrGet](#) over all the leftmost objects to be inserted, then all the right objects to be inserted, yielding tagged relations between [ObjectIds](#). For each inserted relation, the existing relation instances are

removed from those to insert, and the remaining inserted to the relevant [RelationTable](#) with the correct [CommitId](#). The auxiliary tables are now updated and, on success, the views table is updated.

3.10.5 Mutability

As with the memory backend, all mutability except for the availability of views and the default view is hidden from the user. The SQL backend uses commits to manage view mutability.

3.10.6 Pathfinding and Fixed Point Traversal

Pathfinding is implemented by constructing an SQL query to generate right hand side [ObjectIds](#) for a relation given a left hand side [ObjectId](#). This query is used as the search function for the generic pathfinding algorithms. Once paths have been found, their [ObjectIds](#) are looked up in the database to find the full values along each path.

Fixed point traversal and repetitions are done natively in SQL. [FixedPoint](#) and [Upto](#) are done using a recursive CTE, while [Exactly](#) is done by explicitly joining together the required number of repetitions of the sub-relations query.

Recursive CTE example. [Caption](#) = An example of a recursive CTE computing [Upto](#). [FixedPoint](#) would omit the counter variable and the limit.

3.10.7 Object Storage

Objects are stored as in the appropriate [ObjectTable](#) in a manner derived from the [DBObject](#) of each object. Each [DBCCell](#) is converted to appropriate SQL type.

3.11 LMDB Backends

The final family of backends are the LMDB backends. LMDB is a simple, efficient memory mapped file based Key-Value datastore. (<https://symas.com/lmdb/>). It is already the backend for DBMSs ([Source?](#)).

3.11.1 Common

Although I have written several versions of the LMDB backend, there are several overarching similarities.

LMDB API

Terminology LMDB terminology differs slightly from that of more mainstream systems. What would be called a database in SQL is known as an [Env](#) and a table known as a [DBI](#).

Transactions LMDBs transactional model differs from more traditional systems. In short, we need to create a transaction to do reads but not writes. This stems from LMDB delegating as much work as possible to the OS kernel's memory-mapped file system. All writes are immediately written to the database, whereas starting a read transaction gives a view of the `Env` as it was at the start of the transaction, ensuring that concurrent writes don't affect what is read from the database. Writes may be included in transactions to allow atomic get-and-sets. In practice the transactions we actually use are very short lived.

JVM API The LMDB backend uses the `LMDBJava` API <https://github.com/lmdbjava/lmdbjava> to allow access to an LMDB Database from JVM languages (such as Scala). For each `DBI`, we have a byte array to byte array key-value map, much like a `Map[Array[Byte], Array[Byte]]`.

Storage and Keys In order to write to the database using the `LMDBJava` API, we need to convert the key and value to `ByteBuffer`. To generify this process I have introduced two type classes: `Keyable` and `Storable`.

figures for these type-classes `Keyable` is a type class which shows that an object can be converted into an array of bytes representing a component of a key (Can be concatenated with others to make a full key). `Storable` is a more general typeclass that shows that objects can be converted into a `ByteBuffer`. `Storable` objects are typically more complex than objects used as keys, and have a higher marshalling throughput, hence the lower level, faster, `ByteBuffer` API. Typeclass instances exist to allow sets and lists of `Storable` objects to be stored.

Table Structure Similarly to the SQL implementation, there are several tables, some generated based on the `SchemaSescription` and other, control, tables exist regardless. Each table is represented in memory by a subclass of the `impl.lmdb.common.tables.interfaces.LMDBTable` trait, which provides utility methods such as transactional reads and computations.

Import control table from G-drive `Singleton Keys` denote tables with single key.
Import schema table from g-drive

For each `SchemaObject` in the `Schemadescription`, there exists a `RetrievalTable`, `EmptyIndexTable`, and a `ColumnIndexTable` per column in the `SchemaObject`.

Writes Writes occur in a similar way to the SQL implementation, again using commits to manage views. First we lookup all the left and right hand side objects to find their `ObjectIds` (If needed, we create new entries in the retrieval table and update the index tables to insert the objects.), we then look up existing instances of the relations to insert and only insert new ones to the relation table with the new `CommitId`. Upon success, we create a new view and insert it to the views table and the available views table.

3.11.2 Original LMDB Implementation

[impl.lmdb.original](#)

The original LMDB backend implementation executed queries in much the same way as the in-memory backend, the only difference being the flat LMDB table structure. Interpretation occurs by a very similar to function that of the memory backend, instead passing around a list of Commits to search rather than the memory tree. At the end of a query's execution, the resulting [ObjectIds](#) are looked up in the relevant Retrieval table to extract the actual objects and convert them into user objects. Pathfinding and the repetition operators are handled in the same way as the memory backend by calling the relevant generic algorithms.

3.11.3 Batched

[impl.lmdb.fast](#) This is the first variant upon my simple original LMDB backend, it focuses on making small local optimisations rather than significant algorithmic changes.

Read Batching Firstly, the original implementation made suboptimal uses of reads. For example, when extracting objects at the end of query execution, it would lookup each object individually with its own read transaction. This is sped up by batching together reads to separate keys and commits in the same transaction, allowing for a small speed up.

Pre-Caching The result of executing a [FindSingle](#) is constant throughout the whole query, so by lifting out and pre-emptively executing any [FindSingles](#), we speed up execution by removing redundancy. More importantly, this helps enable the next operation.

FindFrom For the pathfinding commands, the transitive queries and [From FindSingle](#) query, were not actually interested in finding all pairs matching a query, or interested in which pairs are related, but actually the reachability function $A \Rightarrow Set[B]$ (this is a kleisli arrow), which takes a node and finds the directly reachable neighbours. Since we don't have to do any bookwork to maintain the left hand side of any relation, this greatly simplifies our search. Hence, in these cases, we interpret the query using a new [FindFrom](#) method, which is an alternative interpreter. For example, joins now become a concatenation of these arrows ([flatMap](#) in Scala parlance), [Id](#) becomes the return function $(x \Rightarrow Set(x))$, [Distinct](#) becomes set subtraction as opposed to calling the [Set.filter](#) function, and, finally, we can make use of the [SimpleFixedPointTraversal](#) algorithms described above. The pre-caching above is useful here, since due to the [flatMap](#)s, the computation of [FindSingles](#) would be repeated many times otherwise.

3.11.4 Common Subexpression Elimination

[impl.lmdb.cse](#) A common performance issue in the previous implementation was that of redundancy and lack of general common subexpression elimination (CSE). When a common sub expression is repeated outside of an [Upto](#) or [Exactly](#) block, then it is repeatedly

recalculated, leading to poor performance. In addition, there is lots of redundancy exposed by the above [FindFrom](#) optimisation. Consider calculation $Rjoin(SjoinT)$. If a node n appears in the results of $S(m)$, $S(m')$ for distinct $(?, m) \triangleleft_{A,B,v} R$, $(?, m') \triangleleft_{A,B,v} R$ (m, m' are right hand results of R), then $T(n)$ is computed redundantly in the overall calculation. These issues of redundancy can be solved using a combination of global CSE and memoization.

The Memoisation Problem When looking at a graph data structure, memoization is harder than when looking to optimise pure functions. To optimise a pure function $f: A \Rightarrow B$, we only need to store a hash table [Map\[A, B\]](#). To compute $f(a)$, we simply look up a in the map. If it is in the map, we return the mapped value else, we call f , and add the result to the map before returning. In a large graph, computing and storing all pairs generated by a query is extremely wasteful, since many of those pairs might never be used (effectively, we would have ignored the left-optimisation). Conversely, memoizing the full left-optimised function, $interpret(Q): Set[A] \Rightarrow Set[(A, B)]$ is also wasteful, since there may be overlap in the sets used as keys. Instead, we memoise the [FindFrom](#) function of a query, [get this findfrom to be verbatim](#) $findFrom(Q): A \Rightarrow Set[B]$, and reconstruct results from the pairs. This allows for the best overlap.

To make the best use of memoization, and to overusing memory by memoizing the same query twice, we need to also apply CSE to group together instances of the same query.

Retrievers A solution to these parameters is the use of an object called a [Retriever](#). This exposes two methods which mirror those from interpreting a query. [Insert retriever trait here from G-Drive](#) For each subquery node in a query (e.g. [And\(Id`A, Or\(R`1, R`2\)\)](#)), we generate a retriever. Retrievers for primitive relations (e.g. [Id`A](#), [Rel\(R\)](#)) are uncached, as LMDB allows for very fast re-computations of these. In the other cases, we use a [CachedRelationRetriever](#), which memorises an underlying lookup function. Using type-enrichment, I have implemented methods such as [join](#), [and](#), [or](#), [exactly](#), [fixedPoint](#) on [RelationRetrievers](#), allowing them to be composed to mirror the query they are generated from. This memoises every node of the query to be executed, yielding a significant reduction of redundancy.

Monadic Compilation To avoid storing multiple [RelationRetrievers](#) for the same subexpression, it is useful to reuse the same retriever for every occurrence of a subquery. Hence when building a retriever for a query, we want to keep track of all subqueries we've seen before. This can be done using the monadic compilation pattern above. The compilation state stores a hashmap of all the subtrees that have already been computed (A bit like when constructing a binary decision diagram in automated theorem proving).

[Compilation state from G-Drive](#) [Compilation state](#) The compilation function checks the memo for precomputed values, and if they are not already computed, recurses to compute subtrees, then composes the subtree results to get a new retriever for the current node, which is added to the compilation state.

3.11.5 Complex Common Subexpression Elimination

`impl.lmdb.fastjoins` Despite these optimisations, there are still sources of redundancy.

Index Building Firstly, the original CSE implementation uses the excessively generic join function to join result sets. This function wastes time by computing an index map on every call. So instead we can change the definition of a `Retriever`'s public functions to mitigate the need to build indices. `new retriever trait` Indeed, the function to join a pair of retrievers becomes a simpler `flatMap` of one map to another. Functions such as the `union` and `intersection` of retrievers also become simpler, (see `impl.lmdb.fastjoins.retrievers.RelationRetriever.RelationRetrieverOps`)

Exactly This implementation also explores other formulations of the `Exactly(n, P)` query. `Exactly` effectively joins together n repetitions of the underlying query in a linear fashion. This makes relatively poor reuse of the join function, since each join has different operand subqueries, so it cannot be memoised. This formulation does make good reuse of `P`, however.

Diagram showing n joins of `P`, linear

As joins over the set of results of query form a `Monoid Proof in appendix or in evaluation?`, we can use associativity to change the order in which the joins are evaluated, aiming for as much reuse as possible. This can be done in a manner similar to binary exponentiation (<http://computingonline.net/computing/article/viewFile/229/204>).

We calculate retrievers for P^{2^i} for each i less than the bitlength of n . We then join the relevant retrievers to get a retriever for P^n . This only uses $O(\log_2(n))$ distinct joins, and makes very good reuse of the join functions. `Diagram showing join structure; Algorithm`

Upto Optimisation We can also re-formulate an `Upto(n, P)` as an `Exactly(n, Or(P, IdA))` `Proof in appendix`, and hence use the same optimisation as above on `Upto`. The same cannot be said for `FixedPoint`, as we do not know statically at what point the underlying relation will converge, so cannot pick an n to compute `Upto(n, P)`.

Chapter 4

Evaluation

4.1 Unit Tests

The correctness of each backend is tested using a suite of 45 unit tests which verify adherence to semantics. These have wide range, testing over all the commands, all the possible ADT nodes, and the correct usage and separation of views. I have been using the regression test model, in that discovering a non-trivial bug, Ive written a test case to target that bug, ensuring that it is not leaked into production again.

4.2 Performance Tests

4.2.1 Introduction

In order to evaluate the effectiveness of the various optimisations to the LMDB backends described above, it was necessary to run performance tests over wide range of queries. The various LMDB backends were tested against each other as a control and particularly against the SQL backend as a standard to beat. The in-memory backend was omitted as initial tests indicated that it runs with approximately the same algorithmic characteristics as the Original LMDB implementation. The LMDB implementation was able to do indexing faster than Scalas tree based hash maps, which are used by the in-memory implementation. Further more, the original LMDB implementation is typically around 3x slower than the batched version for most jobs, so we also omit it for most tests, since most of the algorithms it uses are the same.

4.2.2 Hardware

Tests were run on the oslo machine belonging to Timothy Jones group. Its specifications are shown below. All of the backends ran off of an SSD. [Oslo Specs](#)

4.2.3 Datasets

To evaluate tests on non-trivial examples, I sought to construct datasets over which large queries were feasible.

IMDB The first and most used collection of tests are derived from the IMDB movie database, the most popular 5000 of which were collected from kaggle <https://www.kaggle.com/tmdb/tmdb-movie-metadata>

The CSV data was processed using a python script (in [src/resources/imdb/](#)) into a simplified JSON format. A separate, larger dataset was constructed from tests for another graph database. <https://github.com/arangodb/example-datasets/tree/master/Graphs/IMDB> and manually (in python) converted to the same JSON format. I then wrote a Scala script which reads the JSON and writes the relations to a given database instance.

Slightly different parameters were given when generating each dataset. [Table from GDrive with all the data](#)

The objects in the database are as follows: [Object and relation tables](#)

UFC A second dataset, built from UFC fight data, <https://www.kaggle.com/cformey24/ufc-fights-data-1993-2232016/data> was constructed in the same way to produce one JSON dataset.

[object and relation tables](#)

The [ShorterThan](#) and [LighterThan](#) relations produce extremely sparse graphs which take a long time to converge under transitive closure. [Size of graph.](#)

4.2.4 Test Harness

In order to run tests against each other, I have written a typesafe test harness to run on oslo. This standardises the interface that individual test instances must implement. Each Test must have a [setup](#) method and a [test](#) method which is run on a [DBInstance](#) with the test method indexed by a [TestIndex](#). The test specification must give a maximum index and a mask to avoid running inappropriate backends (For example, those that might take too long on a large test.) The benchmarks that I have run test only the read speed. The time taken to construct the database is not included.

4.2.5 Results

Overall Picture An overall view of the results is that, as might be expected, the SQL implementation is the fastest, with the most aggressively optimised LMDB typically performing the closest in speed to SQL, [calculate precise values](#) though occasionally beating SQLs performance by up to 150%. Due to a high overhead of initiating individual queries, the SQL instance performed appalling on pathfinding queries. The LMDB optimisations worked very effectively where they were hypothesised to, such as queries with repeated joins and those which made use of [Exactly](#) and [Upto](#), yielding orders of magnitude speedup. However, in some tests, the optimisations yielded overheads slowing down performance, especially those queries which make use of [FindFrom](#) rather than finding pairs.

Typically, the CSE optimisations alone do not provide a statistically significant speed up, or even slow down operations, however the optimisations also applied to joins and transitive queries do often show a large speedup.

Pathfinding

Redundancy

Conjunctions and Disjunctions

Tests that involve repetitions

Exactly Test

Exactly Pairs

UptoTest

UptoLarge

JoinSpeed

4.3 Semantics Proofs

4.3.1 Denotational == Operational

4.3.2 Typesafety

4.3.3 Join as a Monoid

Chapter 5

Conclusion